

HPC Benchmark Toolkit

Technical Report

A Comprehensive Framework for Benchmarking
LLM Inference Services on HPC Clusters

Target Platform:	MeluXina HPC Cluster
Container Runtime:	Apptainer/Singularity
Scheduler:	Slurm
Services:	Ollama, vLLM (Single & Distributed)

Contents

1	Introduction	4
1.1	Project Overview	4
1.1.1	Motivation	4
1.1.2	Key Features	4
1.2	Supported Services	4
2	System Architecture	4
2.1	High-Level Architecture	5
2.2	Design Patterns	5
2.2.1	Factory Pattern	5
2.2.2	Template Method Pattern	6
2.3	Component Details	6
2.3.1	Server Managers	6
2.3.2	Workload Controllers	6
2.3.3	Workload Executors	7
3	Communication Architecture	7
3.1	Overview	7
3.2	Control Plane Communication	7
3.3	Data Plane Communication	7
3.4	Metrics Collection Pipeline	8
3.5	Message Sequence Diagram	9
4	Execution Flow	9
4.1	Seven-Phase Execution Model	9
4.2	Detailed Phase Descriptions	10
4.2.1	Phase 1: Initialization	10
4.2.2	Phase 2: Service Deployment	10
4.2.3	Phase 3: Monitoring Setup	10
4.2.4	Phase 4: Client Launch	11
4.2.5	Phase 5: Execution	11
4.2.6	Phase 6: Teardown	11
4.2.7	Phase 7: Report Generation	12
5	Recipe Configuration System	12
5.1	Recipe Structure	12
5.2	Recipe Validation	13
5.3	Parameter Sweeps	13
6	Distributed Benchmarking with Ray	13
6.1	Ray Cluster Architecture	13
6.2	RayClusterManager	14
6.3	Distributed vLLM Configuration	14
7	Monitoring System	15
7.1	How Monitoring Works	15
7.2	Data Collection on HPC	15
7.3	Data Viewing on Your Laptop	16
7.4	Step-by-Step: From HPC to Your Dashboard	16
7.5	What the Dashboards Show	17
7.6	SSH Tunneling for Metrics	17

7.7	SSH Tunneling for Metrics Access	18
8	Benchmarking Implementation	19
8.1	Workload Executor Implementation	19
8.2	Ollama Benchmarking	19
8.3	vLLM Benchmarking	19
8.4	Metrics Collection	19
8.5	Load Patterns	19
9	Logging System	21
9.1	Logging Architecture	21
9.2	BaseLogCollector	21
9.3	TailerLogCollector	21
9.4	Log Categories	21
9.5	Log Format and Storage	21
10	CLI and User Interface	22
10.1	benchmark_cli.py	22
10.2	Orchestrator Arguments	22
10.3	Interactive Recipe Creation	22
11	Slurm Integration	22
11.1	SBATCH Script Generation	22
11.2	Resource Allocation	23
12	Extensibility	23
12.1	Adding a New Service	23
12.2	Custom Metrics	24
13	Deployment and Operations	24
13.1	Prerequisites	24
13.2	Python Dependencies	24
13.3	Container Setup	25
14	Performance Considerations	26
14.1	Bottleneck Analysis	26
14.2	Optimization Recommendations	26
14.3	Scaling Analysis	26
15	Troubleshooting Guide	26
15.1	Common Issues	27
15.2	Debugging Commands	27
16	Project Structure	27
17	Division of Work	28
17.1	Alberto Finardi	28
17.2	Giovanni	29
17.3	Laura	29
17.4	Giulia	30
17.5	Contribution Summary	31
18	Conclusion	31
18.1	Future Work	31

A Port Reference	32
B Environment Variables	32

1 Introduction

1.1 Project Overview

The HPC Benchmark Toolkit is a production-ready framework designed for benchmarking Large Language Model (LLM) inference services on High-Performance Computing (HPC) clusters. The toolkit addresses the critical need for reproducible, scalable, and observable benchmarking in enterprise AI deployments.

1.1.1 Motivation

Modern AI deployments require careful performance characterization before production rollout. Key challenges include:

- **Reproducibility:** Experiments must be repeatable with identical configurations
- **Scalability:** Benchmarks must scale from single-node to multi-node distributed setups
- **Observability:** Real-time metrics are essential for performance analysis
- **HPC Integration:** Seamless integration with Slurm and containerized workloads

1.1.2 Key Features

Feature	Description
Multi-Service Support	Ollama, vLLM (single and distributed), extensible architecture
Distributed Benchmarking	Multi-node server/client orchestration with Ray
Real-Time Monitoring	Prometheus/Grafana integration with live dashboards
Recipe-Driven	YAML configurations for reproducible experiments
HPC-Optimized	Slurm integration, Apptainer container support
Comprehensive Metrics	Latency (p50/p90/p99), throughput, resource utilization

Table 1: Key features of the HPC Benchmark Toolkit

1.2 Supported Services

The toolkit currently supports the following inference services:

Service	Description	Default Port	API Type
Ollama	Local LLM inference server	11434	REST
vLLM	High-throughput LLM serving	8000	OpenAI-compatible
vLLM Distributed	Multi-node vLLM with Ray	8000	OpenAI-compatible
Dummy	Template for custom services	5000	REST

Table 2: Supported inference services

2 System Architecture

2.1 High-Level Architecture

The system follows a modular architecture with clear separation of concerns. Figure 1 illustrates the high-level component structure.

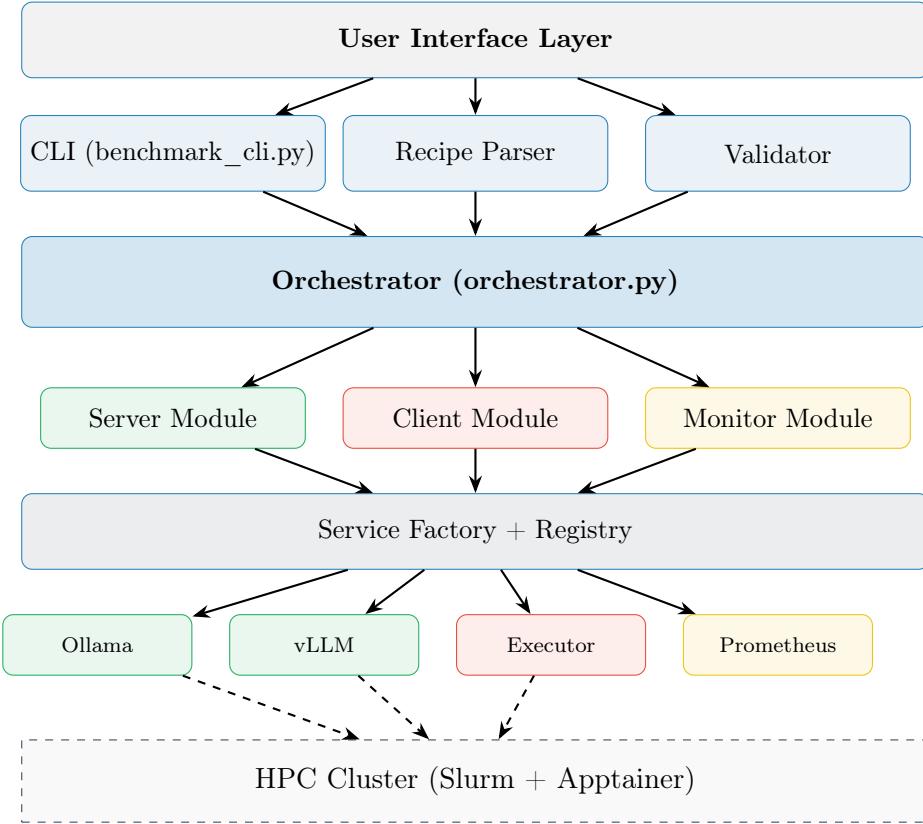


Figure 1: High-level system architecture

2.2 Design Patterns

The toolkit employs several software design patterns to ensure maintainability and extensibility:

2.2.1 Factory Pattern

The **ServiceFactory** class provides dynamic component instantiation based on service type:

```

1  class ServiceFactory:
2      _registry = {}
3
4      @classmethod
5      def register_service(cls, name, server_cls, controller_cls, executor_cls):
6          cls._registry[name] = {
7              'server': server_cls,
8              'controller': controller_cls,
9              'executor': executor_cls
10         }
11
12     @classmethod
13     def create_server_manager(cls, service_name, config):
14         return cls._registry[service_name]['server'](config)
    
```

Listing 1: ServiceFactory implementation pattern

2.2.2 Template Method Pattern

Base classes define algorithm skeletons while subclasses provide specific implementations:

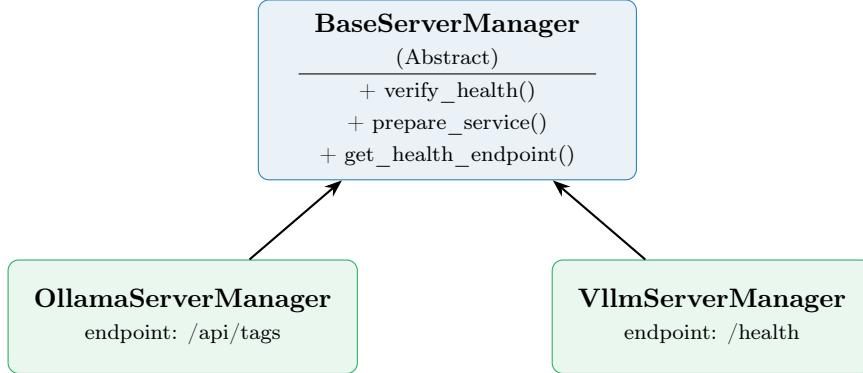


Figure 2: Template Method Pattern in Server Managers

2.3 Component Details

2.3.1 Server Managers

Server managers handle the lifecycle of inference services:

- **Health Checking:** Polls service endpoints to verify readiness
- **Service Preparation:** Loads models, initializes resources
- **Configuration Parsing:** Extracts service-specific settings from recipes

Manager	Health Endpoint	Port	Special Features
OllamaServerManager	/api/tags	11434	Model pulling via /api/pull
VllmServerManager	/health	8000	Ray cluster integration
DummyServerManager	/health	5000	Template implementation

Table 3: Server Manager implementations

2.3.2 Workload Controllers

Controllers coordinate workload execution across client nodes:

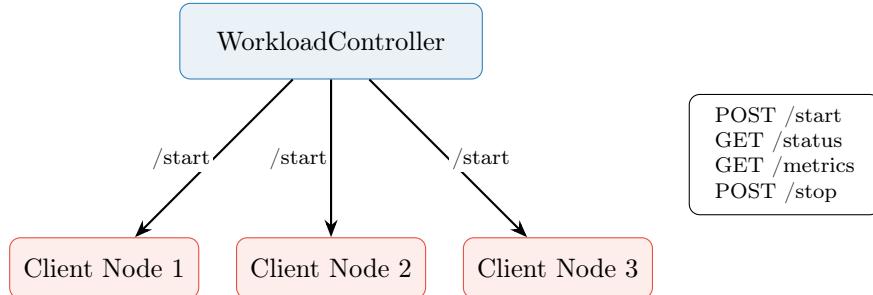


Figure 3: Workload Controller to Client communication

2.3.3 Workload Executors

Executors run on client nodes as Flask servers, executing the actual benchmark workload:

```

1 # Flask endpoints exposed by each executor
2 GET /health           # Check executor status
3 POST /start            # Start workload with JSON config
4 GET /status             # Get current workload status
5 GET /metrics            # Fetch collected metrics
6 GET /metrics/prometheus # Prometheus-compatible format
7 POST /stop              # Stop workload execution

```

Listing 2: Workload Executor REST API

3 Communication Architecture

3.1 Overview

The toolkit employs a hybrid communication architecture combining:

- **REST/HTTP**: Control plane communication between orchestrator and components
- **Service APIs**: Data plane communication between clients and inference servers
- **Prometheus Push/Pull**: Metrics collection and aggregation

3.2 Control Plane Communication

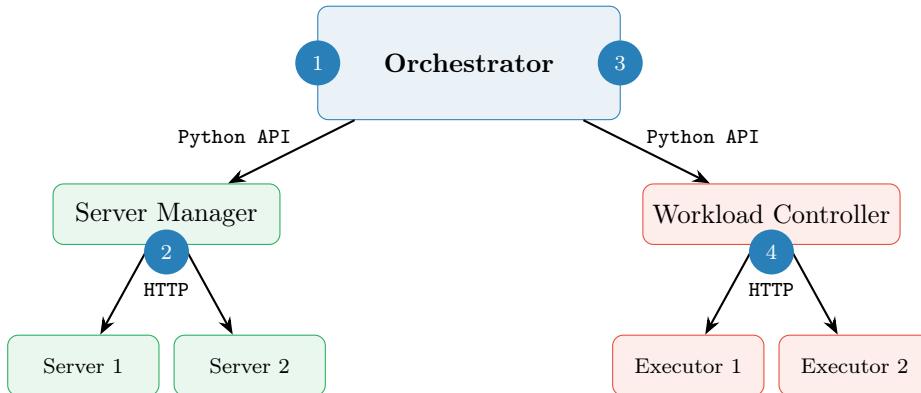


Figure 4: Control plane communication flow

3.3 Data Plane Communication

During benchmark execution, clients send inference requests directly to servers:

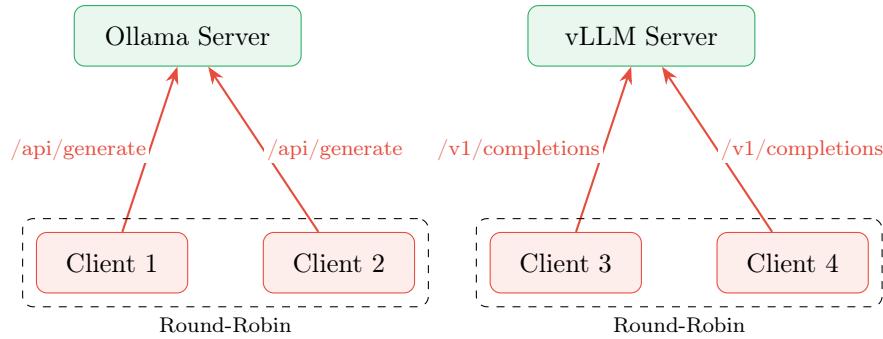


Figure 5: Data plane: Client to Server inference requests

3.4 Metrics Collection Pipeline




Figure 6: Metrics collection pipeline with SSH tunneling

3.5 Message Sequence Diagram

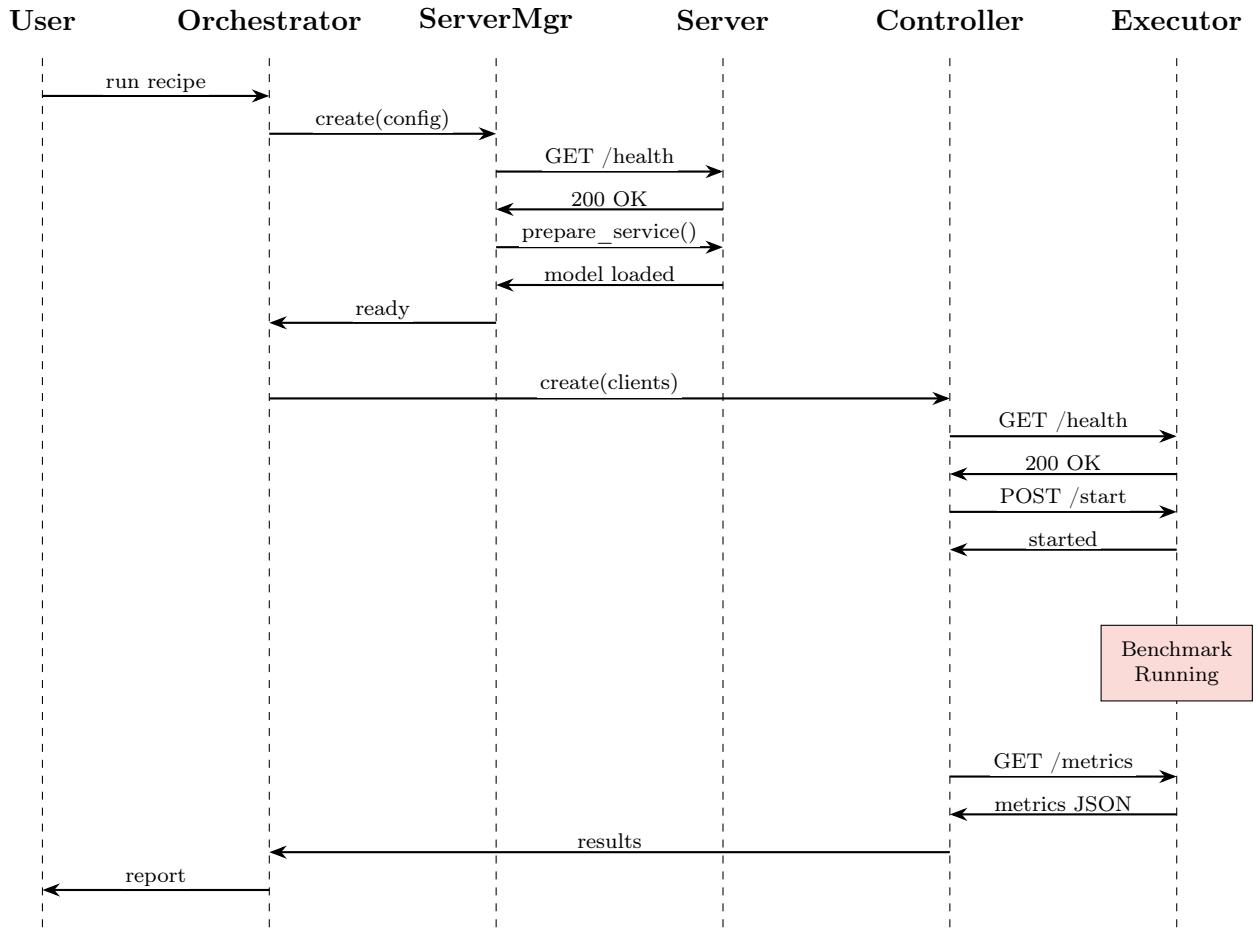


Figure 7: Message sequence for benchmark execution

4 Execution Flow

4.1 Seven-Phase Execution Model

The benchmark execution follows a structured seven-phase model:

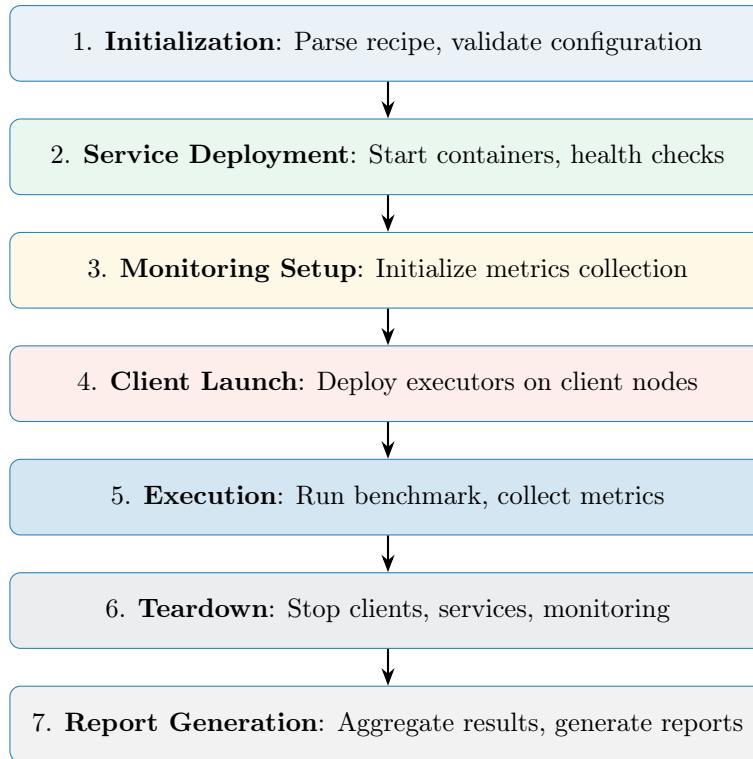


Figure 8: Seven-phase execution model

4.2 Detailed Phase Descriptions

4.2.1 Phase 1: Initialization

1. Load and parse YAML recipe file
2. Validate against JSON schema (`schemas/recipe-format.yaml`)
3. Expand parameter sweeps (Cartesian product)
4. Initialize logging and output directories

4.2.2 Phase 2: Service Deployment

1. Create ServerManager via ServiceFactory
2. Build endpoint list from server node hostnames
3. Poll health check endpoints with configurable timeout
4. Execute service preparation (model loading)

4.2.3 Phase 3: Monitoring Setup

This phase sets up the system to track what's happening during the benchmark. It prepares both the HPC cluster and personal laptop to collect performance data:

1. **Start Monitoring:** Turn on the monitor tool that will watch CPU, GPU, and memory usage
2. **Choose What to Track:** Decide which metrics to collect (CPU usage, GPU usage, memory, etc.)

3. **Connect to Pushgateway:** Set up connection to Prometheus Pushgateway (stores the data)
4. **Create Folders:** Make directories to save the collected data as CSV files
5. **Run Monitor in Background:** Start monitoring that runs continuously and takes a measurement every 1 second by default
6. **Check Grafana Dashboard:** Make sure the dashboard on your laptop is ready to see the data

Main Settings:

- `monitor_interval`: How often to check the system (example: every 1 second)
- `prometheus_push_interval`: How often to send data to storage (example: every 15 seconds)
- `pushgateway_url`: Where to send the data (example: `http://mel2109:9091`)
- `output_file`: File name to save results (example: `benchmark_metrics.csv`)

Parts of the Monitoring System:

- **Prometheus on Your Laptop** (port 9092): Collects data every 15 seconds from the HPC cluster
- **Grafana on Your Laptop** (port 3001): Shows graphs and charts of the performance data
- **Pushgateway on HPC**: A storage box that receives metrics from the cluster
- **SSH Tunnels**: Secret paths that let your laptop see the HPC metrics safely

4.2.4 Phase 4: Client Launch

1. Create WorkloadController via ServiceFactory
2. Verify executor health on all client nodes
3. Distribute workload configuration
4. Initialize thread pools on each executor

4.2.5 Phase 5: Execution

1. Warmup period (configurable duration)
2. Main benchmark execution
3. Continuous metrics collection
4. Real-time Prometheus scraping

4.2.6 Phase 6: Teardown

1. Send stop signals to all executors
2. Collect final metrics
3. Stop monitoring processes
4. Clean up resources

4.2.7 Phase 7: Report Generation

1. Aggregate metrics from all clients
2. Calculate statistics (mean, p50, p90, p99)
3. Generate CSV/JSON output files
4. Create summary report

5 Recipe Configuration System

5.1 Recipe Structure

Recipes are YAML files that declaratively define benchmark configurations:

```

1 scenario: "experiment-name"
2 partition: "gpu"
3 account: "p200981"
4 qos: "default"
5
6 orchestration:
7   mode: "slurm"
8   total_nodes: 5
9   node_allocation:
10    servers:
11      nodes: 2
12    clients:
13      nodes: 2
14      clients_per_node: 10
15    monitors:
16      nodes: 1
17 job_config:
18   time_limit: "02:00:00"
19   exclusive: true
20
21 resources:
22   servers:
23     gpus: 2
24     cpus_per_task: 1
25     mem_gb: 32
26   clients:
27     gpus: 0
28     cpus_per_task: 2
29     mem_gb: 16
30
31 workload:
32   component: "inference"
33   service: "ollama"
34   duration: "2m"
35   warmup: "1m"
36   model: "llama2"
37   clients_per_node: 10
38
39 servers:
40   health_check:
41     enabled: true
42     timeout: 300
43     interval: 5
44     endpoint: "/api/tags"
45   service_config:
46     gpu_layers: 0

```

```

47
48 artifacts:
49   containers_dir: "/path/to/containers/"
50   service:
51     path: "ollama_latest.sif"
52   python:
53     path: "python_3_12_3_v2.sif"
54
55 binds:
56 - "/project/.ollama:/root/.ollama:rw"
57 - "/project/scratch:/scratch:rw"

```

Listing 3: Complete recipe structure

5.2 Recipe Validation

PLACEHOLDER:

Laura's Section Please describe the recipe validation system in detail:

- JSON Schema validation implementation
- Validation rules and error messages
- Custom validators
- Schema versioning
- Interactive validation mode

5.3 Parameter Sweeps

The recipe system supports automatic parameter expansion:

```

1 workload:
2   batch: [1, 4, 8]          # 3 values
3   concurrency: [1, 8, 32]    # 3 values
4   prompt_len: [128, 512]    # 2 values
5   # Total trials: 3 x 3 x 2 = 18

```

Listing 4: Parameter sweep configuration

6 Distributed Benchmarking with Ray

6.1 Ray Cluster Architecture

For distributed vLLM deployments, the toolkit integrates with Ray for tensor and pipeline parallelism:

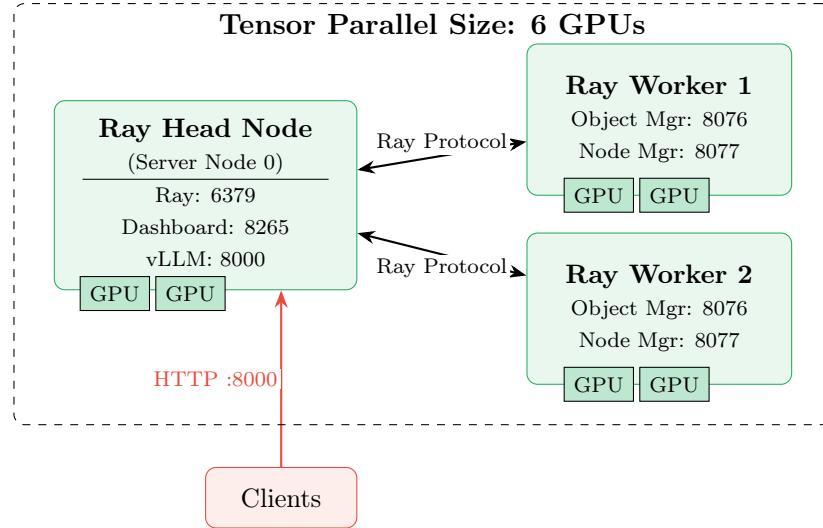


Figure 9: Ray cluster architecture for distributed vLLM

6.2 RayClusterManager

The `RayClusterManager` class handles Ray cluster lifecycle:

```

1  class RayClusterManager:
2      def start_head_node(self, port: int = 6379) -> bool:
3          """Initialize Ray head node with ray start --head"""
4          cmd = f"ray start --head --port={port}"
5          # Execute and verify
6
7      def connect_worker(self, head_address: str) -> bool:
8          """Connect worker to existing Ray cluster"""
9          cmd = f"ray start --address={head_address}"
10         # Execute and verify
11
12     def get_head_ip(self) -> str:
13         """Auto-detect local IP for Ray communication"""
14         # Network interface detection

```

Listing 5: RayClusterManager key methods

6.3 Distributed vLLM Configuration

```

1 servers:
2     service_config:
3         distributed:
4             enabled: true
5             backend: "ray"
6             tensor_parallel_size: 4
7             pipeline_parallel_size: 1
8             ray:
9                 dashboard_port: 8265
10                object_manager_port: 8076
11                node_manager_port: 8077
12                num_cpus_per_node: 4
13                num_gpus_per_node: 2
14                max_model_len: 2048
15                gpu_memory_utilization: 0.7

```

Listing 6: Distributed vLLM recipe configuration

7 Monitoring System

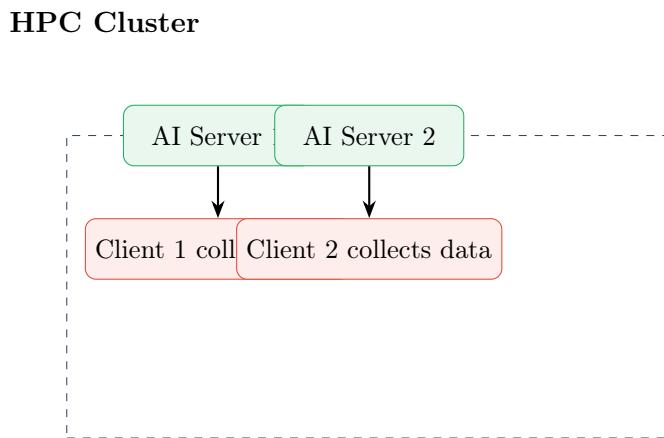
The monitoring system watches what happens during the benchmark and shows you the results. It has two main parts: one on the HPC cluster that collects data, and one on your laptop that shows graphs of this data.

7.1 How Monitoring Works

Simple Overview:

1. Client computers on HPC measure CPU, GPU, and memory usage every second
2. These measurements are stored and sent to a storage server
3. Your laptop connects to the HPC cluster through a secure tunnel
4. Your laptop collects these measurements every 15 seconds
5. Your laptop shows you graphs of all this data in Grafana

7.2 Data Collection on HPC

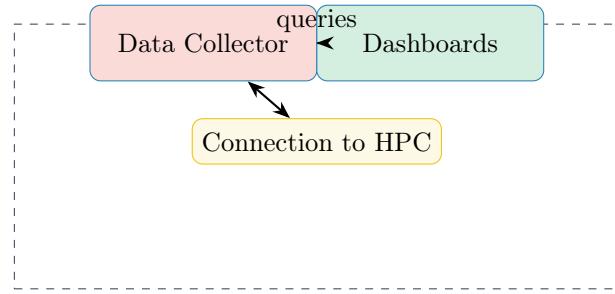


On the HPC cluster, each client computer has a small Python program that:

1. Measures system performance (CPU, GPU, memory) every 1 second
2. Stores these measurements in memory
3. Provides them through a web endpoint on port 6000

7.3 Data Viewing on Your Laptop

Your Laptop



On the laptop, there are two programs:

1. **Prometheus** (on port 9092): A program that collects data from HPC every 15 seconds through the secure tunnel
2. **Grafana** (on port 3001): A program that creates nice graphs from the collected data

7.4 Step-by-Step: From HPC to Your Dashboard

Step 1: HPC Client Collects Data

On client computers on HPC, a Python monitor continuously runs:

```

# Data is available here on HPC (inside the cluster)
http://client-node:6000/metrics/prometheus

# The data looks like:
ollama_throughput_rps 578          # 578 requests per second
ollama_workload_running 1           # 1 = workload is running
ollama_request_latency_seconds 0.015 # Each request takes 15ms

```

Step 2: Create a Tunnel to HPC

Since your laptop cannot directly see HPC, a secure tunnel is created:

```

# Example: Create tunnel from your port 25000 to HPC client-node port 6000
ssh -N -L 25000:client-node:6000 meluxina &

# Now you can see HPC data from your laptop at:
http://localhost:25000/metrics/prometheus

```

Step 3: Prometheus Reads the Data

A program called Prometheus on your laptop reads from this tunnel:

```

# Prometheus configuration file tells it where to look
# It checks every 15 seconds
scrape_configs:
  - job_name: 'ollama'
    static_configs:
      - targets: ['localhost:25000', 'localhost:25001']

```

Step 4: Grafana Shows You the Graphs

Grafana reads the stored data from Prometheus and shows nice graphs:

- A number showing if the workload is running (0 means stopped, 1 means running)
- A line graph showing requests per second over time

- A line graph showing response time over time
- Counters showing total successes and failures

7.5 What the Dashboards Show

Ollama Performance Dashboard

This shows performance of the Ollama AI model:

- **Is it Running?**: Shows 0 (stopped) or 1 (running)
- **Total Requests**: How many AI requests have been made (goes up over time)
- **Speed per Computer**: Shows requests/sec on each computer (line graph)
- **Total Speed**: Shows all requests/sec combined (big number)
- **Response Time**: How long each request takes (line graph in milliseconds)
- **Success vs Failure**: Comparison of successful vs failed requests (bar chart)
- **Error Percentage**: What percent of requests failed (line graph)

vLLM Performance Dashboard

Same metrics of the Ollama but for the vLLM AI model, with extra panels:

- **Tokens per Request**: How many words the AI generates per request
- **Cache Usage**: How full the GPU cache is (as a percentage)

7.6 SSH Tunneling for Metrics

- **Ollama Dashboard**: 13 panels monitoring Ollama workload metrics
 - Workload Status (gauge: 0=idle, 1=running)
 - Requests Total (counter)
 - Throughput by Host (timeseries)
 - Total Throughput (stat panel in RPS)
 - Request Latency (timeseries in ms)
 - Success vs Failed Requests (bar chart)
 - Error Rate Over Time (percentage)
 - Plus additional CPU/GPU/Memory panels
- **vLLM Dashboard**: Same layout with vLLM-specific metrics
 - vllm_workload_running
 - vllm_throughput_rps
 - vllm_request_latency_seconds
 - vllm_tokens_per_request
 - vllm_cache_usage

Example Grafana Queries:

```
# Current workload status
ollama_workload_running

# Throughput aggregated across hosts
sum(ollama_throughput_rps)

# Latency p95
histogram_quantile(0.95, ollama_request_latency_seconds)

# Error rate
100 * (sum(rate(ollama_errors_total[5m])) / sum(rate(ollama_requests_total[5m])))
)
```

7.7 SSH Tunneling for Metrics Access

1. Find client nodes from job output:

```
ssh meluxina "scontrol show job JOBID | grep StdOut"
ssh meluxina "cat /path/to/logs/*.out" | grep "Client nodes"
```

2. Open SSH tunnels:

```
# Ollama clients
ssh -N -L 25000:mel2120:6000 meluxina &
ssh -N -L 25001:mel2148:6000 meluxina &

# vLLM clients
ssh -N -L 25002:mel2142:6000 meluxina &
ssh -N -L 25003:mel2185:6000 meluxina &
```

3. Verify metrics endpoint:

```
curl http://localhost:25000/metrics/prometheus | head -5
```

4. Access Grafana:

```
open http://localhost:3001 # admin / admin
```

8 Benchmarking Implementation

PLACEHOLDER:

Laura's Section Please provide comprehensive documentation of the benchmarking system:

8.1 Workload Executor Implementation

- Thread pool management
- Request generation
- Latency measurement
- Error handling

8.2 Ollama Benchmarking

- HellaSwag dataset integration
- Request format
- Response parsing
- Metrics collection

8.3 vLLM Benchmarking

- OpenAI-compatible API usage
- Streaming vs non-streaming
- Token counting
- Throughput calculation

8.4 Metrics Collection

- Latency percentiles (p50/p90/p99)
- Throughput (requests/second, tokens/second)
- Error rates
- Resource utilization

8.5 Load Patterns

- Constant load
- Poisson distribution
- Burst patterns

Please include:

- Code snippets for key implementations
- Diagrams showing request flow
- Example metrics output
- Performance considerations

9 Logging System

PLACEHOLDER:

Giulia's Section Please provide comprehensive documentation of the logging system:

9.1 Logging Architecture

- Overall logging design
- Log sources and destinations
- Aggregation strategy

9.2 BaseLogCollector

- Abstract interface design
- LogSource dataclass
- Method specifications

9.3 TailerLogCollector

- File tailing implementation
- Remote node log collection
- Log aggregation

9.4 Log Categories

- Application logs
- System logs (Slurm)
- Benchmark logs
- Infrastructure logs

9.5 Log Format and Storage

- Structured logging (JSON)
- Timestamps and correlation IDs
- Storage organization
- Retention policies

Please include:

- TikZ diagrams for log flow
- Code snippets for key implementations
- Example log entries
- Integration with other components

10 CLI and User Interface

10.1 benchmark_cli.py

The main CLI provides three primary commands:

Command	Arguments	Description
list	-	Display all available recipes with details
create	-	Interactive wizard for recipe creation
run	-recipe PATH	Deploy and run a benchmark recipe

Table 4: CLI commands

10.2 Orchestrator Arguments

```
python3 orchestrator.py \
--server-nodes NODE [NODE ...]           # Required: Server hostnames
--client-nodes NODE [NODE ...]           # Required: Client hostnames
--workload-config-file PATH              # Required: Recipe file
[--server-port PORT]                    # Default: 11434/8000
[--client-port PORT]                    # Default: 5000
[--timeout SECONDS]                     # Default: 600
[--enable-monitoring]                   # Enable metrics
[--pushgateway-node NODE]               # For Prometheus
[--monitor-interval SECONDS]           # Default: 5
[--monitor-output PATH]                 # Output file
```

Listing 7: Orchestrator command-line arguments

10.3 Interactive Recipe Creation

The CLI guides users through recipe creation:

1. Service selection (Ollama, vLLM, vLLM Distributed)
2. Scenario configuration (name, partition, account)
3. Node allocation (servers, clients, monitors)
4. Resource requirements (GPUs, CPUs, memory)
5. Workload parameters (model, duration, clients)
6. Container paths and bind mounts

11 Slurm Integration

11.1 SBATCH Script Generation

The toolkit generates Slurm batch scripts from recipes:

```
#!/bin/bash
#SBATCH --job-name=ollama-benchmark
#SBATCH --partition=gpu
#SBATCH --account=p200981
#SBATCH --nodes=5
#SBATCH --time=02:00:00
```

```
#SBATCH --exclusive

# Load modules
module load Apptainer

# Get node list
NODES=($(scontrol show hostnames $SLURM_JOB_NODELIST))
SERVER_NODES="${NODES[0]} ${NODES[1]}"
CLIENT_NODES="${NODES[2]} ${NODES[3]}"
ORCHESTRATOR_NODE="${NODES[4]}"

# Start servers
for node in $SERVER_NODES; do
    srun --nodes=1 --nodelist=$node \
        apptainer run --nv ollama.sif &
done

# Wait for servers
sleep 30

# Start client executors
for node in $CLIENT_NODES; do
    srun --nodes=1 --nodelist=$node \
        apptainer exec python.sif \
        python3 executor.py --port 6000 &
done

# Run orchestrator
python3 orchestrator.py \
    --server-nodes $SERVER_NODES \
    --client-nodes $CLIENT_NODES \
    --workload-config-file recipe.yaml
```

Listing 8: Generated SBATCH script structure

11.2 Resource Allocation

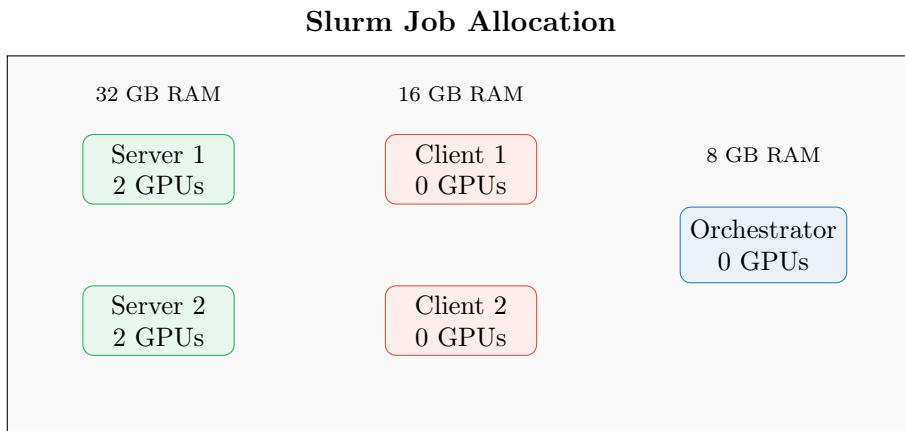


Figure 10: Slurm resource allocation example

12 Extensibility

12.1 Adding a New Service

To add a new inference service, implement four components:

1. **Server Manager:** Handle service lifecycle
2. **Workload Controller:** Coordinate clients
3. **Workload Executor:** Execute benchmarks
4. **Service Registration:** Register with factory

Files to Create/Modify

```
servers/myservice_server_manager.py
controller/myservice_workload_controller.py
executor/myservice_workload_executor.py
service_registry.py (update)
```

12.2 Custom Metrics

Extend the Monitor class for custom metrics:

```
1 from prometheus_client import Gauge
2
3 class CustomMonitor(Monitor):
4     def __init__(self, *args, **kwargs):
5         super().__init__(*args, **kwargs)
6         self.custom_metric = Gauge(
7             'custom_metric',
8             'Description of custom metric'
9         )
10
11     def collect_custom(self):
12         value = self._get_custom_value()
13         self.custom_metric.set(value)
```

Listing 9: Custom metrics extension

13 Deployment and Operations

13.1 Prerequisites

Component	Requirement	Version
Python	Runtime	3.6+
Slurm	Job scheduler	Any
Apptainer	Container runtime	1.0+
Docker	Local monitoring	20.10+

Table 5: System prerequisites

13.2 Python Dependencies

```
pip install flask requests psutil prometheus_client pyyaml
```

Listing 10: Required Python packages

13.3 Container Setup

```
module load Apptainer

# Pull Ollama container
apptainer pull docker://ollama/ollama:latest

# Pull vLLM container
apptainer pull docker://vllm/vllm-openai:latest

# Pull Python container for clients
apptainer pull docker://python:3.12.3-slim
```

Listing 11: Building containers on MeluXina

14 Performance Considerations

PLACEHOLDER:

Performance Analysis - Pending Laura's Benchmarking Results This section will be completed once Laura's benchmarking work is finalized. It will include:

14.1 Bottleneck Analysis

- Latency breakdown for LLM inference
- GPU compute vs memory transfer analysis
- Network overhead measurements
- Tokenization performance impact

14.2 Optimization Recommendations

- Warmup period requirements
- Batch size tuning guidelines
- Tensor parallelism configuration
- Memory utilization optimization
- Client concurrency tuning

14.3 Scaling Analysis

- Single-node vs multi-node performance
- Strong scaling efficiency
- Weak scaling characteristics
- Resource utilization patterns

Please provide:

- Benchmark results from Ollama and vLLM tests
- Performance charts and graphs
- Latency distributions (p50/p90/p99)
- Throughput measurements
- Resource utilization data
- Optimization recommendations based on findings

15 Troubleshooting Guide

15.1 Common Issues

Issue	Symptom	Solution
Server health check fails	Timeout during startup	Verify port accessibility, check container logs
Client cannot connect	Connection refused	Check executor is running, verify port
No metrics in Grafana	Empty dashboard	Verify SSH tunnel, check Prometheus targets
Ray cluster fails	Workers not connecting	Check network ports, verify head IP

Table 6: Common issues and solutions

15.2 Debugging Commands

```
# Check server health
curl http://server-node:11434/api/tags

# Verify executor
curl http://client-node:6000/health

# Check Prometheus targets
curl http://localhost:9092/api/v1/targets

# View Ray cluster status
ssh head-node "ray status"
```

Listing 12: Useful debugging commands

16 Project Structure

```
hpc-benchmark-toolkit/
++ src/
|   +-- benchmark/
|   |   +-- orchestrator.py
|   |   +-- service_factory.py
|   |   +-- service_registry.py
|   |   +-- servers/
|   |   |   +-- base_server_manager.py
|   |   |   +-- ollama_server_manager.py
|   |   |   +-- vllm_server_manager.py
|   |   |   +-- ray_cluster_manager.py
|   |   +-- workload/
|   |   |   +-- controller/
|   |   |   +-- executor/
|   |   +-- logging/
|   +-- benchmark_cli.py
|   +-- monitor/
|       +-- monitor.py
+-- monitoring/
|   +-- docker-compose.yml
|   +-- prometheus.yml
|   +-- grafana/
+-- schemas/
|   +-- recipe-format.yaml
+-- docs/
```

```
+-- diagrams/
```

Listing 13: Complete project structure

17 Division of Work

This section documents the contributions of each team member to the project.

17.1 Alberto Finardi

Role: System Architect, Infrastructure, and Core Implementation

Contributions:

- **Core Infrastructure**
 - Designed and implemented the overall system architecture
 - Created the Service Factory pattern for extensibility
 - Built the modular component structure
 - Implemented service registry for dynamic component loading
- **Server Management**
 - Implemented `BaseServerManager` abstract class
 - Developed `OllamaServerManager` with health checks and model pulling
 - Developed `V11mServerManager` with distributed support
 - Created `RayClusterManager` for distributed vLLM deployments
- **Workload Controllers**
 - Designed `BaseWorkloadController` interface
 - Implemented service-specific controllers (Ollama, vLLM)
 - Developed HTTP-based client coordination protocol
- **Workload Executors**
 - Created `BaseWorkloadExecutor` Flask server framework
 - Implemented REST API endpoints for workload management
 - Developed thread pool management for concurrent requests
 - Built metrics collection infrastructure
- **Benchmarking Framework**
 - Implemented core benchmarking logic
 - Developed request generation and load patterns
 - Created latency measurement and metrics aggregation
 - Built Ollama and vLLM specific benchmark implementations
- **CLI Development**
 - Developed `benchmark_cli.py` with three main commands
 - Implemented interactive recipe creation wizard

- Created deployment and job submission logic
- Built recipe listing and management features

- **Orchestration**

- Implemented main orchestrator logic (`orchestrator.py`)
- Developed seven-phase execution model
- Created Slurm sbatch script generation
- Built node allocation and resource management

- **Basic Logging Infrastructure**

- Set up initial logging framework
- Implemented basic log output and formatting
- Created log directory structure

- **Documentation**

- Wrote comprehensive README.md
- Created API Reference documentation
- Developed Developer Guide
- Authored this technical report

- **Integration**

- Integrated all team contributions
- Coordinated component interfaces
- Performed system testing and debugging

17.2 Giovanni

Role: Monitoring Integration

Contributions:

- **Prometheus Integration**

- Configured Prometheus for metrics collection
- Set up Docker container for Prometheus server

- **Grafana Dashboards**

- Created Grafana dashboards for Ollama and vLLM metrics visualization
- Configured Docker container for Grafana

17.3 Laura

Role: Recipe Validation and Benchmark Execution

Contributions:

- **Recipe Validation System**

- Designed JSON Schema for recipe validation (`schemas/recipe-format.yaml`)
- Implemented validation logic and error handling

- Created custom validators for service-specific configurations
- Developed user-friendly error message formatting
- Built schema versioning support

- **Benchmark Execution**

- Ran and validated Ollama benchmarks on MeluXina
- Ran and validated vLLM benchmarks (single and distributed)
- Tested parameter sweep configurations
- Verified metrics collection and reporting

- **Documentation**

- Recipe validation section of this report
- Benchmarking implementation section of this report

17.4 Giulia

Role: Logging System

Contributions:

- **Logging Architecture**

- Designed overall logging architecture
- Defined log sources and destinations
- Created aggregation strategy for distributed logs

- **BaseLogCollector**

- Implemented abstract interface for log collection
- Designed `LogSource` dataclass
- Defined method specifications for collectors

- **TailerLogCollector**

- Implemented file tailing for real-time log collection
- Built remote node log collection via SSH
- Created log aggregation from multiple sources

- **Log Management**

- Implemented structured logging format (JSON)
- Added timestamps and correlation IDs
- Organized storage structure
- Defined retention policies

- **Documentation**

- Logging system section of this report

17.5 Contribution Summary

Team Member	Primary Area	Key Deliverables
Alberto Finardi	Infrastructure & Core	Orchestrator, Server Managers, Controllers, Executors, Benchmarks, CLI, Basic Logging, Documentation
Giovanni	Monitoring	Prometheus/Grafana integration, Dashboards
Laura	Validation	Recipe validation, Benchmark execution
Giulia	Logging	Log collectors, Aggregation

Table 7: Team contribution summary

18 Conclusion

The HPC Benchmark Toolkit provides a comprehensive solution for benchmarking LLM inference services on HPC clusters. Key achievements include:

- **Modular Architecture:** Clean separation of concerns with factory pattern
- **Multi-Service Support:** Ollama, vLLM (single and distributed)
- **Real-Time Monitoring:** Full Prometheus/Grafana integration
- **Reproducibility:** YAML-based recipe system
- **Extensibility:** Easy addition of new services
- **HPC Integration:** Native Slurm and Apptainer support

18.1 Future Work

Potential enhancements include:

- Support for additional inference services (TensorRT, Triton)
- Kubernetes orchestration mode
- Automated performance regression testing
- Interactive web dashboard
- Multi-cluster support

A Port Reference

Port	Service	Description
11434	Ollama	Default Ollama API
8000	vLLM	Default vLLM API
5000/6000	Executor	Workload executor Flask
9091	Pushgateway	Prometheus Pushgateway
9092	Prometheus	Prometheus server
3001	Grafana	Grafana dashboard
6379	Ray	Ray cluster communication
8265	Ray Dashboard	Ray monitoring
8076	Ray Object Manager	Ray object store
8077	Ray Node Manager	Ray node management

Table 8: Complete port reference

B Environment Variables

Variable	Description	Default
PYTHONPATH	Python module path	–
MLUX_USER	MeluXina username	–
MLUX_ACCOUNT	MeluXina project account	–
MLUX_KEY	Path to SSH key	~/.ssh/id_ed25519_mlux

Table 9: Environment variables